

## About the Data

1. Day\_of\_week: The day of the week when the accident occurred.
2. Age\_band\_of\_driver: The age group or band of the driver involved in the accident.
3. Sex\_of\_driver: The gender of the driver involved in the accident.
4. Educational\_level: The educational level of the driver involved in the accident.
5. Vehicle\_driver\_relation: Relationship of the driver with the vehicle (e.g., owner, renter).
6. Driving\_experience: Experience level of the driver in terms of years.
7. Type\_of\_vehicle: Type of vehicle involved in the accident (e.g., car, truck, motorcycle).
8. Owner\_of\_vehicle: Ownership status of the vehicle (e.g., self-owned, company-owned).
9. Service\_year\_of\_vehicle: Number of years the vehicle has been in service.
10. Vehicle\_movement: Movement or action of the vehicle before or during the accident.
11. Casualty\_class: Classification of the casualty (e.g., driver, passenger, pedestrian).
12. Sex\_of\_casualty: Gender of the casualty involved in the accident.
13. Age\_band\_of\_casualty: Age group or band of the casualty involved in the accident.
14. Casualty\_severity: Severity of the casualty (e.g., minor injury, serious injury, fatality).
15. Work\_of\_casualty: Occupation or work status of the casualty.
16. Fitness\_of\_casualty: Fitness status of the casualty.
17. Pedestrian\_movement: Movement of any pedestrians involved in the accident.
18. Cause\_of\_accident: The cause or reason for the accident.
19. Accident\_severity: Severity of the accident itself.

## Importing the libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [2]: #Reading the dataset
df = pd.read_csv('RTA Dataset.csv')
```

```
In [3]: df.head(5)
```

Out[3]:

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_driver_relatic
0	17:02:00	Monday	18-30	Male	Above high school	Employee
1	17:02:00	Monday	31-50	Male	Junior high school	Employee
2	17:02:00	Monday	18-30	Male	Junior high school	Employee
3	1:06:00	Sunday	18-30	Male	Junior high school	Employee
4	1:06:00	Sunday	18-30	Male	Junior high school	Employee

5 rows × 32 columns



```
In [4]: #Checking the information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12316 entries, 0 to 12315
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Time                                  12316 non-null  object
1   Day_of_week                          12316 non-null  object
2   Age_band_of_driver                   12316 non-null  object
3   Sex_of_driver                        12316 non-null  object
4   Educational_level                     11575 non-null  object
5   Vehicle_driver_relation              11737 non-null  object
6   Driving_experience                   11487 non-null  object
7   Type_of_vehicle                      11366 non-null  object
8   Owner_of_vehicle                     11834 non-null  object
9   Service_year_of_vehicle              8388 non-null   object
10  Defect_of_vehicle                    7889 non-null   object
11  Area_accident_occured                12077 non-null  object
12  Lanes_or_Medians                     11931 non-null  object
13  Road_allignment                      12174 non-null  object
14  Types_of_Junction                   11429 non-null  object
15  Road_surface_type                    12144 non-null  object
16  Road_surface_conditions              12316 non-null  object
17  Light_conditions                     12316 non-null  object
18  Weather_conditions                  12316 non-null  object
19  Type_of_collision                    12161 non-null  object
20  Number_of_vehicles_involved           12316 non-null  int64
21  Number_of_casualties                  12316 non-null  int64
22  Vehicle_movement                     12008 non-null  object
23  Casualty_class                       12316 non-null  object
24  Sex_of_casualty                      12316 non-null  object
25  Age_band_of_casualty                  12316 non-null  object
26  Casualty_severity                    12316 non-null  object
27  Work_of_casualty                      9118 non-null   object
28  Fitness_of_casualty                  9681 non-null   object
29  Pedestrian_movement                  12316 non-null  object
30  Cause_of_accident                    12316 non-null  object
31  Accident_severity                     12316 non-null  object
dtypes: int64(2), object(30)
memory usage: 3.0+ MB
```

Their are "29" object columns and "2" numerical(Integer) Columns

```
In [5]: #checking for null values
df.isnull().sum()
```

```
Out[5]: Time                                0
Day_of_week                                0
Age_band_of_driver                         0
Sex_of_driver                             0
Educational_level                         741
Vehicle_driver_relation                    579
Driving_experience                         829
Type_of_vehicle                           950
Owner_of_vehicle                          482
Service_year_of_vehicle                   3928
Defect_of_vehicle                         4427
Area_accident_occured                     239
Lanes_or_Medians                          385
Road_allignment                           142
Types_of_Junction                         887
Road_surface_type                         172
Road_surface_conditions                   0
Light_conditions                          0
Weather_conditions                        0
Type_of_collision                         155
Number_of_vehicles_involved               0
Number_of_casualties                     0
Vehicle_movement                         308
Casualty_class                           0
Sex_of_casualty                          0
Age_band_of_casualty                     0
Casualty_severity                        0
Work_of_casualty                         3198
Fitness_of_casualty                      2635
Pedestrian_movement                      0
Cause_of_accident                        0
Accident_severity                        0
dtype: int64
```

```
In [6]: print(31-15) #Their are 16 columns with null values present
```

16

```
In [7]: #describing the data for numerical data
df.describe()
```

Out[7]:

	Number_of_vehicles_involved	Number_of_casualties
count	12316.000000	12316.000000
mean	2.040679	1.548149
std	0.688790	1.007179
min	1.000000	1.000000
25%	2.000000	1.000000
50%	2.000000	1.000000
75%	2.000000	2.000000
max	7.000000	8.000000

```
In [8]: # Describing the data including categorical columns
df.describe(include='object')
```


Out[8]:

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_driver_
count	12316	12316	12316	12316	11575	
unique	1074	7	5	3	7	
top	15:30:00	Friday	18-30	Male	Junior high school	Er
freq	120	2041	4271	11437	7619	

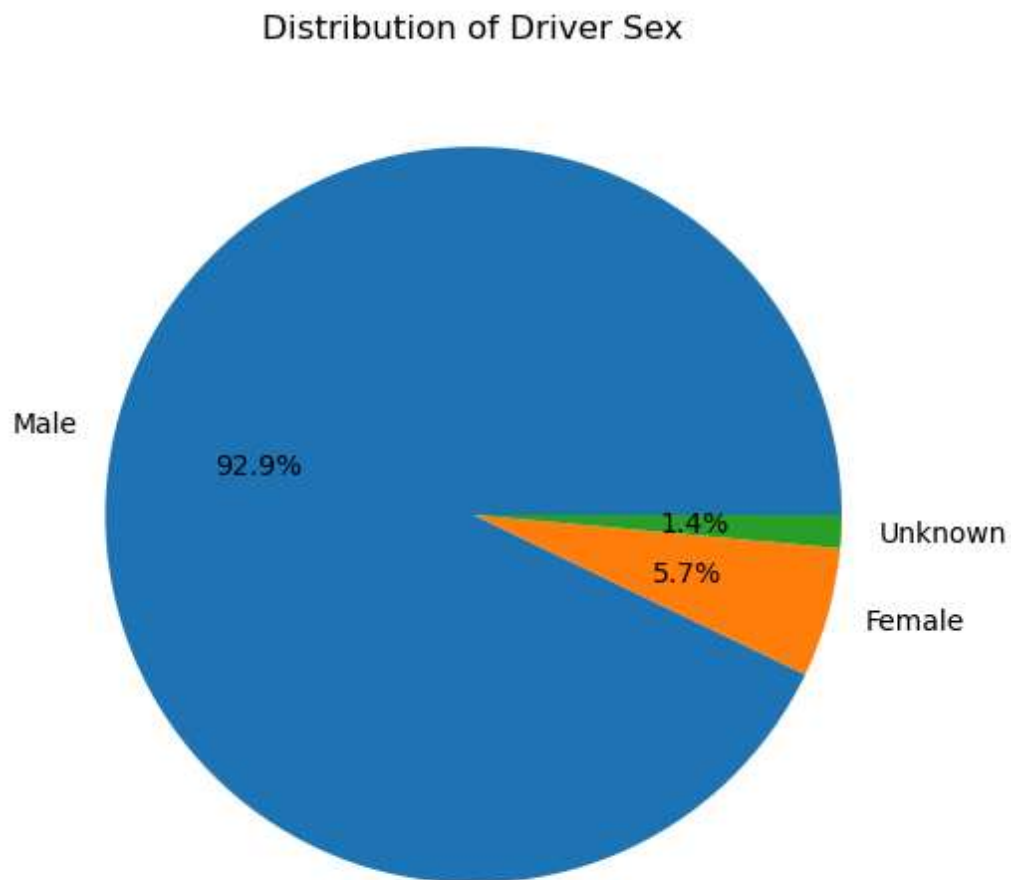
4 rows × 30 columns



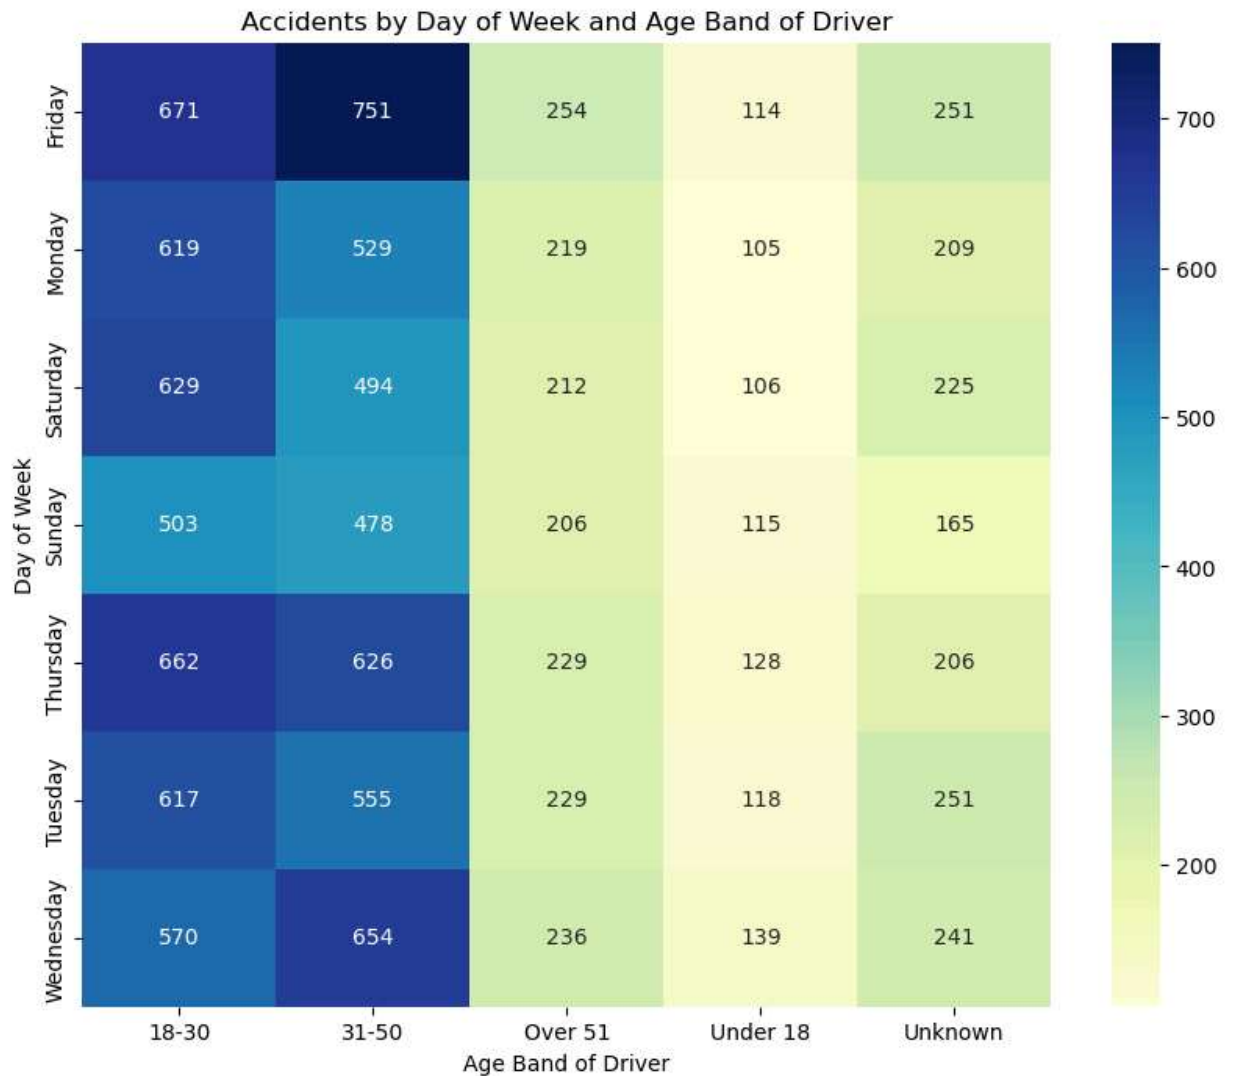
```
In [9]: sns.countplot(x='Day_of_week', data=df)
plt.title('Accidents by Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Number of Accidents')
plt.show()
```



```
In [10]: plt.figure(figsize=(8, 6))
df['Sex_of_driver'].value_counts().plot(kind='pie', autopct='%1.1f%%')
plt.title('Distribution of Driver Sex')
plt.ylabel('')
plt.show()
```



```
In [11]: plt.figure(figsize=(10, 8))
heatmap_data = df.groupby(['Day_of_week', 'Age_band_of_driver']).size().unstack(f
sns.heatmap(heatmap_data, cmap='YlGnBu', annot=True, fmt='d')
plt.title('Accidents by Day of Week and Age Band of Driver')
plt.xlabel('Age Band of Driver')
plt.ylabel('Day of Week')
plt.show()
```



```
In [12]: #Checking for unique/ different type of values present in "Educational_Level" column
df['Educational_level'].unique()
```

```
Out[12]: array(['Above high school', 'Junior high school', nan,
               'Elementary school', 'High school', 'Unknown', 'Illiterate',
               'Writing & reading'], dtype=object)
```

```
In [13]: #Calculating the mode of "Educational_Level"
ELM=df['Educational_level'].mode()
ELM
```

```
Out[13]: 0    Junior high school
Name: Educational_level, dtype: object
```



```
In [14]: #Calculating the mode of "Vehicle_driver_relation"
VDRM = df['Vehicle_driver_relation'].mode()
VDRM
```

```
Out[14]: 0    Employee
Name: Vehicle_driver_relation, dtype: object
```

```
In [15]: #Calculating the mode of "Driving_experience"
DE = df['Driving_experience'].mode()
DE
```

```
Out[15]: 0    5-10yr
Name: Driving_experience, dtype: object
```

```
In [16]: print("The percentage of data that is null:")
df.isnull().sum()/len(df)*100
```

The percentage of data that is null:

```
Out[16]: Time                                0.000000
Day_of_week                                0.000000
Age_band_of_driver                         0.000000
Sex_of_driver                             0.000000
Educational_level                         6.016564
Vehicle_driver_relation                    4.701202
Driving_experience                         6.731082
Type_of_vehicle                           7.713543
Owner_of_vehicle                          3.913608
Service_year_of_vehicle                   31.893472
Defect_of_vehicle                         35.945112
Area_accident_occured                     1.940565
Lanes_or_Medians                          3.126015
Road_allignment                           1.152972
Types_of_Junction                         7.202014
Road_surface_type                         1.396557
Road_surface_conditions                   0.000000
Light_conditions                          0.000000
Weather_conditions                        0.000000
Type_of_collision                         1.258525
Number_of_vehicles_involved                0.000000
Number_of_casualties                      0.000000
Vehicle_movement                          2.500812
Casualty_class                            0.000000
Sex_of_casualty                           0.000000
Age_band_of_casualty                      0.000000
Casualty_severity                         0.000000
Work_of_casualty                          25.966223
Fitness_of_casualty                       21.394933
Pedestrian_movement                       0.000000
Cause_of_accident                         0.000000
Accident_severity                         0.000000
dtype: float64
```

```
In [17]: # created an function to replace the null values with mode value
def null_value_treatment(col):
    for i in df:
        if df[i].dtypes=='object':
            df[i].fillna(df[i].mode()[0],inplace=True)
        else:
            df[i].fillna(df[i].median(),inplace=True)
```

```
In [18]: for i in df:
        null_value_treatment(i)
```

```
In [19]: df.isnull().sum()/len(df)*100
```

```
Out[19]: Time                                0.0
         Day_of_week                        0.0
         Age_band_of_driver                 0.0
         Sex_of_driver                      0.0
         Educational_level                  0.0
         Vehicle_driver_relation            0.0
         Driving_experience                  0.0
         Type_of_vehicle                    0.0
         Owner_of_vehicle                   0.0
         Service_year_of_vehicle            0.0
         Defect_of_vehicle                  0.0
         Area_accident_occured              0.0
         Lanes_or_Medians                   0.0
         Road_allignment                    0.0
         Types_of_Junction                  0.0
         Road_surface_type                  0.0
         Road_surface_conditions            0.0
         Light_conditions                   0.0
         Weather_conditions                 0.0
         Type_of_collision                  0.0
         Number_of_vehicles_involved        0.0
         Number_of_casualties               0.0
         Vehicle_movement                   0.0
         Casualty_class                     0.0
         Sex_of_casualty                    0.0
         Age_band_of_casualty               0.0
         Casualty_severity                  0.0
         Work_of_casualty                   0.0
         Fitness_of_casualty                0.0
         Pedestrian_movement                0.0
         Cause_of_accident                  0.0
         Accident_severity                  0.0
         dtype: float64
```

```
In [20]: df.sample(10)
```

```
Out[20]:
```

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_driver_re
<b>7716</b>	11:10:00	Thursday	Unknown	Male	Junior high school	
<b>11313</b>	17:40:00	Sunday	Over 51	Male	Junior high school	
<b>36</b>	20:30:00	Friday	18-30	Male	Above high school	Em
<b>5185</b>	18:11:00	Thursday	18-30	Male	Junior high school	
<b>2133</b>	1:32:00	Monday	31-50	Male	Junior high school	Em
<b>4016</b>	8:56:00	Saturday	Under 18	Male	High school	Em
<b>526</b>	13:28:00	Monday	18-30	Male	Junior high school	
<b>4410</b>	16:25:00	Monday	18-30	Male	Unknown	Em
<b>7672</b>	6:50:00	Monday	Under 18	Male	Junior high school	Em
<b>7827</b>	12:51:00	Wednesday	Over 51	Male	Junior high school	Em

10 rows × 32 columns



After filling the null values with mode, there are some columns that have "na" and "unknown" present.

```
In [21]: df['Age_band_of_casualty'].unique()
```

```
Out[21]: array(['na', '31-50', '18-30', 'Under 18', 'Over 51', '5'], dtype=object)
```

```
In [22]: df['Service_year_of_vehicle'].unique()
```

```
Out[22]: array(['Above 10yr', '5-10yrs', 'Unknown', '1-2yr', '2-5yrs', 'Below 1yr'],
              dtype=object)
```

```
In [23]: df['Service_year_of_vehicle'].mode()[0]
```

```
Out[23]: 'Unknown'
```

```
In [24]: df['Age_band_of_casualty'].mode()[0]
```

```
Out[24]: 'na'
```

***As we can see that the mode value of that column is null***

```
In [25]: #Checking the correlation
df.corr()
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_16032\18339900.py:2: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
df.corr()
```

Out[25]:

	Number_of_vehicles_involved	Number_of_casualties
Number_of_vehicles_involved	1.000000	0.213427
Number_of_casualties	0.213427	1.000000

***As we can see that "Number\_of\_vehicles\_involved" has correlation with "Number\_of\_casualties" of 0.2134. A correlation of 0.213 is not particularly strong, but it's not necessarily "bad."***

```
In [26]: df.skew()
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_16032\1665899112.py:1: FutureWarning: The default value of numeric\_only in DataFrame.skew is deprecated. In a future version, it will default to False. In addition, specifying 'numeric\_only=None' is deprecated. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
df.skew()
```

Out[26]:

Number_of_vehicles_involved	1.323454
Number_of_casualties	2.344769
dtype:	float64

***A skewness value of 1.3234, 2.344 indicates that the data is positively skewed (right-skewed) and -3.833 indicates that the data is left skewed***

```
In [27]: # Converting the categorical data into numerical data
LE = LabelEncoder()
```

```
In [28]: def Categorical_numerical(col):
          for i in df:
              df[col] = LE.fit_transform(df[col])
```

```
In [29]: for i in df:
          Categorical_numerical(i)
```

```
In [30]: df.head(4)
```

```
Out[30]:
```

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_driver_relation
0	420	1	0	1	0	0
1	420	1	1	1	4	0
2	420	1	0	1	4	0
3	594	3	0	1	4	0

4 rows × 7 columns



```
In [31]: #treating the outlier with zscore
from scipy.stats import zscore
def outlier_treatment(col):
    zscore1=(abs(zscore(df[col])))
    outlier=zscore1>+3
    median=df[col].median()
    df.loc[outlier,col]=median
```

```
In [32]: for i in df:
    outlier_treatment(i)
```

## Training the Model

```
In [33]: X = df.drop('Accident_severity',axis=1)
```

```
In [34]: y = df.Accident_severity
```

## Balancing the data using smote

```
In [35]: from imblearn.over_sampling import SMOTE
smote = SMOTE()
X_resampled, y_resampled = smote.fit_resample(X, y)
df_resampled = pd.concat([pd.DataFrame(X_resampled), pd.DataFrame(y_resampled)],
```

```
In [36]: # Splitting the data into training and testing
from sklearn.model_selection import train_test_split
```

```
In [37]: x_train,x_test,y_train,y_test = train_test_split(X,y,test_size=0.8,random_state=0)
```

## Logistic Regression

```
In [38]: from sklearn.linear_model import LogisticRegression
LR=LogisticRegression()
```

```
In [39]: LR.fit(x_train,y_train)
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

```
Out[39]: LogisticRegression
LogisticRegression()
```

```
In [40]: pred = LR.predict(x_test)
```

## Evaluating the Model

```
In [41]: accuracy_score(y_test,pred)*100
```

```
Out[41]: 85.86217395717041
```

```
In [43]: predict = LR.predict(x_test)
```

```
In [44]: confusion_matrix(y_test,predict)
```

```
Out[44]: array([[ 0, 1393],
               [ 0, 8460]], dtype=int64)
```

## DecisionTreeClassifier()

```
In [45]: #Decision Tree classifier
DTC = DecisionTreeClassifier()
```

```
In [46]: DTC.fit(x_train,y_train)
```

```
Out[46]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [47]: y_pred = DTC.predict(x_test)
```

```
In [48]: y_pred
```

```
Out[48]: array([2, 1, 2, ..., 2, 2, 2], dtype=int64)
```

## Evaluating the model

```
In [49]: accuracy_score(y_test,y_pred)*100
```

```
Out[49]: 75.55059372779864
```

```
In [50]: f1_score(y_test,y_pred,average='weighted')*100
```

```
Out[50]: 76.20773383900621
```

```
In [51]: confusion_matrix(y_test,y_pred)
```

```
Out[51]: array([[ 286, 1107],  
               [1302, 7158]], dtype=int64)
```

## K NEIGHBORS CLASSIFIER

```
In [52]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [53]: KNN = KNeighborsClassifier(n_neighbors=25)
```

```
In [54]: KNN.fit(x_train,y_train)
```

```
Out[54]: ▾ KNeighborsClassifier  
KNeighborsClassifier(n_neighbors=25)
```

```
In [55]: knn_pred=KNN.predict(x_test)
```

```
In [56]: y_test
```

```
Out[56]: 8348    2
          5386    2
          1783    2
          7077    2
          5437    2
          ..
          9012    2
          886     2
          5878    2
          5876    2
          6558    2
          Name: Accident_severity, Length: 9853, dtype: int64
```

## Evaluating the Model

```
In [57]: accuracy_score(y_test,knn_pred)*100
```

```
Out[57]: 85.86217395717041
```

```
In [58]: f1_score(y_test,knn_pred,average='weighted')*100
```

```
Out[58]: 79.33096616367189
```

```
In [59]: confusion_matrix(y_test,knn_pred)
```

```
Out[59]: array([[ 0, 1393],
                [ 0, 8460]], dtype=int64)
```

```
In [60]: print("The true predictions are:")
          8460+0
```

The true predictions are:

```
Out[60]: 8460
```

```
In [61]: print("The False predictions are:")
          1393+0
```

The False predictions are:

```
Out[61]: 1393
```

```
In [62]: KNN.score(x_train,y_train)
```

```
Out[62]: 0.857896873731222
```

```
In [63]: KNN.score(x_test,y_test)
```

```
Out[63]: 0.8586217395717041
```



# Random Forest Classifier

```
In [64]: from sklearn.ensemble import RandomForestClassifier
```

```
In [65]: RFC = RandomForestClassifier()  
RFC.fit(x_train,y_train)
```

```
Out[65]: 

▼ RandomForestClassifier  
RandomForestClassifier()


```

```
In [66]: model_pred = RFC.predict(x_test)
```

```
In [67]: y_test
```

```
Out[67]: 8348      2  
5386      2  
1783      2  
7077      2  
5437      2  
      ..  
9012      2  
886       2  
5878      2  
5876      2  
6558      2  
Name: Accident_severity, Length: 9853, dtype: int64
```

## Evaluating the model

```
In [68]: accuracy_score(y_test,model_pred)*100
```

```
Out[68]: 85.87232315030955
```

```
In [69]: confusion_matrix(y_test,model_pred)
```

```
Out[69]: array([[ 5, 1388],  
               [ 4, 8456]], dtype=int64)
```

```
In [70]: print("The True predictions are:")  
(4+845)
```

The True predictions are:

```
Out[70]: 849
```

```
In [71]: print("The False predictions are:")  
2+1389
```

The False predictions are:

```
Out[71]: 1391
```

```
In [72]: precision = precision_score(y_test, model_pred)  
recall = recall_score(y_test, model_pred)  
  
# Print the results  
print(f"Precision: {precision:.2f}")  
print(f"Recall: {recall:.2f}")
```

Precision: 0.56  
Recall: 0.00

```
In [73]: f1_score(y_test,model_pred)*100
```

```
Out[73]: 0.7132667617689015
```

## Conclusion

```
1. Logistic Regression Model:  
Algorithm Used: Logistic Regression  
Performance Metrics:  
Accuracy: 85.8621%  
F1-score: 76.32%  
  
2. Random Forest Model:  
Algorithm Used: Random Forest  
Performance Metrics:  
Accuracy: 85.88%  
F1-score: 0.57%  
  
3. K-nearest Neighbour:  
Algorithm Used: KNN Classifier  
Performance Metrics:  
Accuracy: 85.86%  
F1-score: 79.33%  
  
4. Decision Tree  
Algorithm Used: Decsison tree Classifier  
Performance Metrics:  
Accuracy: 75.69%  
F1-score: 76.32%
```

The analysis concludes with the following results:

Logistic Regression achieved an accuracy of approximately 85.86%. Decision Tree Classifier had an accuracy of around 75.69%. K-Nearest Neighbors (KNN) achieved an accuracy of 85.86%. Random Forest Classifier had an accuracy of approximately 85.88%. The models were evaluated based on accuracy, F1-score, and confusion matrices.

Further fine-tuning and feature engineering may improve model performance.

In [ ]: